

Design Note — Lexical vs. Semantic Retrieval for News Recommendation

CS4.406 Information Retrieval & Extraction — Assignment 1

1 What We Built

News recommendation differs structurally from catalog domains like retail or streaming: item relevance half-life is hours, not months, so item cold-start is the *default* state, not an edge case (Wu et al. 2020; Kruse et al. 2024) — motivating content-based candidate retrieval, lexical (BM25) and semantic (embeddings), over collaborative filtering.

This project implements both methods, plus one shared evaluation harness, over MIND and EB-NeRD under a unified schema and temporal-split protocol (ADR-001, ADR-002), comparing methods and datasets on identical footing rather than via method-specific pipelines. Both were pushed through to real Codabench leaderboard submissions on both competitions (§3.5).

2 Choices We Made

Data & splits. Official train/validation splits are used as-is (7-day windows) — the only feasible option given data budgets (ADR-001). A unified three-table schema (articles, impressions, user_history) with mandatory core + dataset-specific optional fields lets one retrieval and evaluation code path run against both datasets (ADR-002).

Lexical retrieval. BM25Okapi over unweighted concatenated title+abstract query text (full click history, stopwords removed), scored as a precomputed sparse weight matrix rather than `rank_bm25`'s naive `get_scores()` — the latter measured at ~10hr projected vs. ~80s actual at MINDsmall-dev scale, a real infeasibility, not a style preference (ADR-005, ADR-006).

Semantic retrieval. One embedding model computed over both datasets (rejecting EB-NeRD's provided embeddings + a separate MIND model, to keep the single-code-path comparison valid) — `paraphrase-multilingual-MiniLM-L12-v2`, chosen over `multilingual-e5-small` after a same-category vs. different-category cosine-similarity discrimination check: MiniLM's gap (0.069, 2.4x ratio) is real and usable; e5-small's `passage:-`prefixed convention compresses nearly all pairs into a narrow 0.75–0.78 band (gap 0.017, ratio 1.02x) — e5's asymmetric query/passage objective doesn't fit this project's symmetric user-profile-vs-article use case (ADR-008). ANN backend is brute-force cosine similarity — 0.99ms/query at MINDsmall scale, 2.44ms/query at MINDlarge's 255,990 users — FAISS stays unjustified at this scale (ADR-008 + Addendum; revisited at 10x scale, §5). User representation is mean-pooled, L2-renormalized embeddings of full click history; zero-history users get an explicit all-zero vector, mirroring BM25's own cold-start reporting posture.

Rejecting EB-NeRD's provided embeddings (above) was argued, not measured, until a follow-up (ADR-008 Addendum, 2026-08-19, same eval code) found the provided artifact wins recall@K and every Q4 accuracy metric by a CI-clear margin (AUC 0.5453 vs. 0.5430) but loses Diversity@10 — a real trade-off, EB-NeRD-only, not a blanket verdict (§6).

Evaluation. Two distinct questions are measured, not conflated: *recall@K* (does retrieval surface the clicked article out of the *whole* corpus, $K \in \{50, 100, 200\}$) and *Q4's ranking metrics* — AUC, MRR, nDCG@5, nDCG@10, intra-list diversity, novelty, catalog coverage — over candidates *already listed in each impression*. Ties break via a deterministic per-impression-seeded pseudo-random permutation (cold users produce exact ties); $K = 10$ for diversity/novelty. Every metric reports 95% bootstrap CIs ($n = 2000$), sliced warm/cold (history length ≥ 5 = warm), except coverage, a point estimate only — set-union statistics are mechanically biased under with-replacement resampling (ADR-007).

3 Observations

3.1 Recall@K — candidate generation over the whole corpus

Dataset (split)	Method	recall@50	recall@100	recall@200
MIND-small dev, overall	BM25	0.73%	1.50%	2.62%
MIND-small dev, overall	Embed	0.90%	1.60%	2.78%
MIND-small dev, warm	BM25	0.73%	1.54%	2.73%
MIND-small dev, warm	Embed	0.93%	1.67%	2.92%
MIND-small dev, cold	BM25	0.72%	1.19%	1.78%
MIND-small dev, cold	Embed	0.65%	1.10%	1.81%
EB-NeRD-demo val (warm only*)	BM25	1.01%	2.13%	4.02%
EB-NeRD-demo val (warm only*)	Embed	0.32%	0.94%	2.57%
EB-NeRD-small val (warm only*)	BM25	0.72%	1.44%	2.77%
EB-NeRD-small val	Embed	0.14%	0.43%	1.21%

*EB-NeRD’s active-user-filtered bundles have zero cold-start users (min history length = 5) by construction. Source: `experiments/{bm25,embed}_{mind,ebnerd,ebnerd_small}_2026-08-10/`. Not run at MINDlarge scale — flagged in §6, not approximated.

3.2 Q4 ranking metrics — AUC / MRR / nDCG (bootstrap 95% CI)

Dataset	Method	AUC overall (CI)	AUC warm	AUC cold
MIND-small dev	BM25	0.569 (0.567–0.571)	0.577	0.524
MIND-small dev	Embed	0.634 (0.632–0.636)	0.644	0.574
MINDlarge dev	BM25	0.570 (0.569–0.571)	0.578	0.522
MINDlarge dev	Embed	0.634 (0.633–0.634)	0.643	0.575
EB-NeRD-demo val	BM25	0.533 (0.528–0.537)	— (=overall)	n/a
EB-NeRD-demo val	Embed	0.544 (0.539–0.549)	—	n/a
EB-NeRD-small val	BM25	0.529 (0.527–0.530)	—	n/a
EB-NeRD-small val	Embed	0.543 (0.541–0.544)	—	n/a

On MIND (both scales), embeddings’ AUC CI does not overlap BM25’s on either cohort. The warm/cold gap is *not* similar between methods: BM25’s is $\Delta 0.052$ – 0.056 , embeddings’ is $\Delta 0.068$ – 0.070 — embeddings lift warm AUC by ~ 0.065 – 0.067 over BM25 but cold by only ~ 0.050 – 0.053 , widening the relative warm/cold disparity rather than closing it.

On EB-NeRD, embeddings’ AUC CI also does not overlap BM25’s — a modest (~ 0.011 – 0.014) edge — despite BM25 winning recall@K by 1.6–2.3x on the *same* corpus (§3.1); recall@K asks whether the click is anywhere in the whole catalog, AUC asks how well the method ranks EB-NeRD’s already-curated, shorter candidate lists (median 9–12 vs. MIND’s 23) — see §6 for the leading explanation.

MRR/nDCG@5/nDCG@10: embeddings lead BM25 by a similar margin on MIND (e.g. MINDlarge nDCG@10: 0.390 vs. 0.350) and a smaller, mixed margin on EB-NeRD (e.g. EB-NeRD-small nDCG@10: 0.459 vs. 0.454).

3.3 Diversity / novelty / coverage (K=10, point estimate for coverage)

Dataset	Method	Diversity@10	Novelty@10	Coverage@10
MIND-small dev	BM25	0.837	16.30	0.083
MIND-small dev	Embed	0.827	16.15	0.078
MINDlarge dev	BM25	0.836	18.09	0.067
MINDlarge dev	Embed	0.828	17.76	0.065
EB-NeRD-small val	BM25	0.795	17.17	0.206
EB-NeRD-small val	Embed	0.789	17.19	0.205

BM25 and embeddings sit within ~ 1 point of each other on every beyond-accuracy metric, on every corpus — accuracy differs meaningfully, beyond-accuracy behavior barely does (n up to 376,471 impressions, so not automatically noise, but not the large effect AUC shows).

3.4 System cost

BM25 at MINDlarge scale: 1.70s index build, 23.5MB sparse weight matrix, ~ 9.8 min projected full retrieval for 255,990 users (ADR-006 Addendum). Embeddings: 264.4s to encode+cache all 72,023 articles, ~ 10.4 min projected retrieval (ADR-008 Addendum). Both stay local on an 8GB machine at this scale — only EB-NeRD’s blind test split (13.5M impressions, no local ground truth) required

Kaggle.

3.5 Leaderboard results (Q5)

Competition	Method	Score	Submission ID	Date
MIND (competitions/13967)	Embed (MiniLM)	0.6195	886468	2026-08-12 13:01
MIND (competitions/13967)	Cohort-gated combiner	0.6192	896696	2026-08-22
EB-NeRD (competitions/2469)	Embed (MiniLM)	0.5404	888045	2026-08-13 23:08
EB-NeRD (competitions/2469)	Contrastive vector	0.5404	896072	2026-08-21 13:01

Both used the embedding method — the clear winner on local dev-set AUC for both datasets (§3.2). EB-NeRD’s first score (0.5404) is close to local validation AUC on `ebnerd_small` (0.5430) — a coherence check, not a claim the two are defined identically.

Second MIND submission — cohort-gated combiner, essentially flat. Local validation predicted a real, CI-clear win at every stage: +0.0027 on MINDsmall-dev (95% CI +0.0018 to +0.0034), +0.0019 on a MINDlarge-dev re-verification (95% CI +0.0015 to +0.0023, ADR-010’s Addendum). The real result: **0.6192 vs. 886468’s 0.6195, a -0.0003 difference** — essentially flat, not the win either local check predicted. No pipeline defect found (line count, malformed-permutation, and all 32 N89741-affected-impression checks all clean). Most likely a genuine population difference — MIND’s dev/test splits are disjoint calendar weeks (ADR-001) — so the edge was built around MINDlarge-dev-specific properties, not a population-independent effect. Same category of finding as the EB-NeRD result below, now observed twice.

Second EB-NeRD submission — contrastive vector, real test set. Both submissions round to the same Score (0.5404) — checked before trusting it: the two `predictions.txt` differ in checksum, and 99.48% of impressions rank differently for the same ID, confirming independent runs, not a duplicate upload. The per-submission detail view (grouped by date, “50% of the testset”) resolves the tie:

	06-01	06-02	06-03	06-04	06-05	06-06	06-07	06-08	MEAN
MiniLM AUC	0.5374	0.5422	0.5755	0.5444	0.5296	0.5330	0.5316	0.5240	0.5397
Contrastive AUC	0.5597	0.5533	0.5504	0.5226	0.5314	0.5489	0.5231	0.5323	0.5402

Contrastive leads 5/8 days and the mean, but thinly (+0.0005, sign flips three times day-to-day). MRR/nDCG@5/nDCG@10 means agree in direction (0.3553 vs. 0.3452; 0.3898 vs. 0.3825; 0.4695 vs. 0.4615), but this AUC margin is far smaller than the **~0.0023 CI-clear gap** ADR-008’s Addendum measured locally (0.5453 vs. 0.5430) — real but small, not a dramatically larger win.

3.6 MIND candidate search: eight other approaches, before G

Screened on MINDsmall-dev vs. the baseline (0.634, §3.2) before committing to G; only G’s routing won CI-clear overall, only F’s combiner won CI-clear on the cold cohort. Full evidence: ADR-010 (D–H)/ADR-011 (I).

Cand.	Method	Result
A	Entity embeddings (TransE)	0.553, loss
B	50/50 BM25+embedding hybrid	0.626, loss
C	Recency-weighted embedding query	0.627, loss
D	Symbolic category/entity overlap	0.613, loss
E	Train-popularity only	0.532, loss
F	7-feature LR combiner	0.626 loss; 0.593 cold win
H1	LogisticRegression, balanced	0.632, loss
H2	HistGradientBoosting	0.599, loss
I	Attention re-ranker, 5 epochs	0.623, loss
I-long	Same, 30 epochs	0.621, worse than I
I-pop	I-long + raw popularity feature	0.513, inconclusive

No untuned signal beats the baseline outright; real wins (F, G) combine signals around `log_popularity` — G’s didn’t hold at real scale (above).

4 Anti-Gaming and Leakage (Q9)

Serving-time-unavailable features — measured, not just avoided. Novelty’s popularity signal comes from the **train split only**, Laplace-smoothed, never the evaluation split (ADR-007). EB-NeRD’s raw `articles.parquet` also carries three lifetime aggregate fields — `total_inviews`, `total_pageviews`, `total_read_time` — unavailable at serving time; `src/datasets/ebnerd.py` never reads them (grep-verified). `scripts/run_leakage_ablation.py` measures what including them would buy: an untuned 50/50 blend of the deployed embedding score with a min-max-normalized leak signal, on `ebnerd_small`

validation (ADR-009):

Metric	Without leaky features	With leaky features (50/50 blend)	CI's overlap?
AUC	0.5430 (0.5415–0.5445)	0.5662 (0.5646–0.5678)	No
MRR	0.3437 (0.3421–0.3452)	0.3638 (0.3623–0.3653)	No
nDCG@5	0.3804 (0.3784–0.3823)	0.4034 (0.4016–0.4051)	No
nDCG@10	0.4591 (0.4574–0.4607)	0.4783 (0.4767–0.4799)	No

Every metric moves by a CI-clear margin from exposure to these fields alone (AUC +0.023 absolute, $\sim 4.3\%$), untuned. $\sim 48\%$ of articles were median-imputed, not zero-filled, so this gap is plausibly a conservative underestimate — direct, measured evidence for excluding these fields entirely (ADR-009).

Leakage-boundary test. `tests/integration/test_leakage.py` asserts no user’s history click occurs at or after their own earliest impression time. **Exists and passes for EB-NeRD** (demo train + validation splits). **Explicitly skip-marked for MIND:** MIND provides only a static per-user article-ID list, no per-click timestamps, so the temporal boundary can’t be independently verified from the data. An acknowledged gap, not an oversight — MIND’s temporal safety instead rests on the dataset’s own documented split construction (ADR-001).

5 Where It Breaks at $10\times$

Every real scaling failure so far was invisible at the previously-exercised scale and appeared at a 5–14x jump, not a 2x one — a pattern, not a one-off.

MINDsmall $\rightarrow$ **MINDlarge** ($\sim 5\times$ users: 50,000 $\rightarrow$ 255,990; $\sim 14\times$ impression-candidate rows) hung in a per-token Python loop never exercised at MINDsmall’s counts; the first fix projected to **$\sim 27\text{GB}$** (measured, not estimated) from one Python object per repeated ID string, resolved via categorical dtype (12.5x reduction, $\sim 2.2\text{GB}$); a second bottleneck (`series.map()` on categorical data, $> 1\text{hr}$, zero progress) was found only by profiling the stuck process. **MINDlarge_test (2.37M) $\rightarrow$ ebnerd_testset (13.5M, 5.7x further)** hit a third, structurally identical failure: a converter materializing a full `list[dict]` before writing projected to $\sim 16\text{GB}$, invisible at the smaller scale — fixed by a generator consuming one row at a time.

Projecting one more 10x from the largest verified scale (MINDlarge: 255,990 users/72,023 articles; EB-NeRD: 13.5M impressions) is inference, not measurement. (1) Retrieval time scales roughly linearly with users — 10x alone projects to ~ 98 min (BM25) / ~ 104 min (embeddings), feasible locally. (2) Brute-force cosine is $O(\text{users} \times \text{articles})$ — if the corpus also grows 10x, that’s 100x compute: ~ 10.4 min $\rightarrow$ roughly **17 hours**, forcing a revisit of ADR-008’s “FAISS unjustified” call. (3) Every bug so far shared one shape — a per-row-object or full-materialization pattern past an unexercised scale; nothing guarantees another path is safe at 10x further.

6 Limitations and Open Questions

- **True (zero-history) cold-start is a shared ceiling**, not solved by either method — MIND’s cold-cohort recall@200 is nearly identical (BM25 1.78% vs. embed 1.81%).
- **EB-NeRD’s cold cohort is structurally empty** in both demo and small (min history length = 5, by construction), so no warm/cold comparison exists for EB-NeRD; `ebnerd_large` is unverified, and recall@K (§3.1) was never run at MINDlarge scale.
- **The Q9 leakage-boundary test is real but partial** — passes for EB-NeRD, skip-marked for MIND (no per-click timestamps; §4). The Q9 ablation’s magnitude is specific to an untuned 50/50 blend — a real, CI-clear lower bound, not the maximum a calibrated leaky model could extract.
- **Two choices rest on a proxy check, not a gold standard:** the encoder uses a category-based discrimination proxy, not a curated near-duplicate set (ADR-008); the hand-built Danish+English stopword list is unvalidated against an NLP resource (ADR-005).
- **The EB-NeRD recall-vs-AUC disagreement isn’t fully explained** (§3.2): BM25 wins whole-corpus recall 1.6–2.3x, yet embeddings edge the curated candidate-list AUC — plausibly the diffuse mean-pooled query, not isolated.
- **Local CI-clear wins haven’t transferred to real blind-test results, twice:** MIND’s combiner ($+0.0027$ local $\rightarrow -0.0003$ real) and EB-NeRD’s contrastive vector ($+0.0023 \rightarrow +0.0005$) — both likely population-specific.